

Lecture 19

Last time

- ▷ ADMM
- ▷ Examples

Today

- ▷ Preconditioning
- ▷ PDHG
- ▷ Primal-dual guarantee

Preconditioning operators

We go back to a template we considered for Fenchel duality

$$(P) \inf_{x \in E} f(x) + g(Ax)$$

linear $f: E \rightarrow \mathbb{R}$
closed, convex, proper

We could apply ADMM, but last time we saw that this would imply solving a linear system involving A , which might be prohibitively expensive!

Instead, we take an alternative path. Recall that its dual was

$$(D) \sup_{y \in Y} -f^*(-A^*y) - g^*(y). \quad \begin{matrix} \text{changed} \\ \text{a sign} \end{matrix}$$

Moreover, a pair (\bar{x}, \bar{y}) are primal-dual solutions iff

$$-A^*\bar{y} \in \partial f(x) \text{ and } \bar{y} \in \partial g(A\bar{x}),$$

$$\Leftrightarrow -A^*\bar{y} \in \partial f(x) \text{ and } A\bar{x} \in \partial g^*(\bar{y}),$$

which means, the solution is a

zero

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} \in \underbrace{\left(\begin{bmatrix} 0 & A^* \\ -A & 0 \end{bmatrix} + \begin{bmatrix} \partial f \\ \partial g^* \end{bmatrix} \right)}_I \begin{bmatrix} \bar{x} \\ \bar{y} \end{bmatrix}.$$

We could aim to apply the KM

iteration by considering R_T

$$\sum_{k+1} \left\{ \begin{bmatrix} x_{k+1} \\ y_{k+1} \end{bmatrix} \leftarrow R_T \begin{bmatrix} x_k \\ y_k \end{bmatrix} \right\}$$

or equivalently

$$\begin{bmatrix} x_k - x_{k+1} \\ y_k - y_{k+1} \end{bmatrix} \in \begin{bmatrix} A^* y_{k+1} + \partial f(x_{k+1}) \\ -A x_{k+1} + \partial g^*(y_{k+1}) \end{bmatrix}$$

This algorithm is mostly conceptual as it requires evaluating a resolvent depending on both primal and dual variables.

Instead we use the following simple, but remarkably useful observation.

Proposition: Suppose $T: E \rightrightarrows E$ is

a monotone operator with a zero. Consider a sequence defined via

$$(A) \quad H(z_k - z_{k+1}) \in \alpha^T z_{k+1} \quad \text{with } \alpha > 0$$

with $H: E \rightarrow E$ a linear, self-adjoint, positive definite. Then, the sequence converges $z_k \rightarrow z^*$ to a zero of T .

Proof: The map H has a square root, call it W , with $H = W \circ W$. Make the change of variable $w \leftarrow Wz$. Then

$$W \circ W(z_k - z_{k+1}) \in \alpha^T z_{k+1}$$



$$W(w_k - w_{k+1}) \in \alpha^T W^{-1} z_{k+1}$$

$$(\omega_k - \omega_{k+1}) \in \alpha W^{-1} T W^{-1} \omega_{k+1}.$$

Then, it is easy to see that S is monotone:

monotonicity of T $\langle S(w) - S(w'), w - w' \rangle$
 $= \langle T(L^{-1}w) - T(L^{-1}w'), L^{-1}(w - w') \rangle \geq 0.$

Then, by the KM iteration we have $\omega_k \rightarrow \omega^*$ with $Sw^* = 0$ and $z_k = W^{-1}\omega_k \rightarrow W^{-1}\omega^* = z^*$ with

$$0 = Sw^* = W^{-1}T W^{-1}\omega^* = W^{-1}Tz^*$$

$$0 \Leftrightarrow W0 = Tz^*. \quad \square$$

Thus, if we find a good H to make the algorithm (A) implementa-

ble.

Primal-Dual Hybrid Gradient

In 2011, Chambolle and Pock came up with a simple H

$$H = \begin{pmatrix} \frac{1}{\tau} I & -A^* \\ -A & \frac{1}{s} I \end{pmatrix} \quad \text{for } \tau, s > 0$$

Claim: H is positive definite as long as $\tau s \|A\|_{op}^2 < 1$.

This H induces the following update

$$\begin{bmatrix} \frac{1}{\tau} I & -A^* \\ -A & \frac{1}{s} I \end{bmatrix} \begin{bmatrix} x_k - x_{k+1} \\ y_k - y_{k+1} \end{bmatrix} \in \begin{bmatrix} A^* y_{k+1} + \partial f(x_{k+1}) \\ -A x_{k+1} + \partial g^*(y_{k+1}) \end{bmatrix}$$

equivalently

$$\begin{aligned} \frac{1}{\tau} (x_k - x_{k+1}) - A^* (y_k - y_{k+1}) &\in A^* y_{k+1} + \partial f(x_{k+1}) \\ -A (x_k - x_{k+1}) + \frac{1}{s} (y_k - y_{k+1}) &\in -A x_{k+1} + \partial g^*(y_{k+1}) \end{aligned}$$

rearranging

$$x_k - \tau A^* y_k \in x_{k+1} + \tau \partial f(x_{k+1})$$

$$y_k + s A (2x_{k+1} - x_k) \in y_{k+1} + s \partial g(x_{k+1})$$

or equivalently

$$x_{k+1} \leftarrow \text{prox}_{\tau f}(x_k - \tau A^* y_k)$$

$$y_{k+1} \leftarrow \text{prox}_{sg^*}(y_k + s A (2x_{k+1} - x_k)).$$

Remarkably each step is now implementable, provided that we can compute the prox of f and g^* .

Example

If we take $f = \langle c, \cdot \rangle + \tau_{R_+^n}(\cdot)$

$g = \tau_{\{x \mid x = b\}}(\cdot)$. Then, we obtain

$$\text{prox}_{\tau f}(x) = \text{proj}_{R_+^n}(x - \tau c)$$

$$\begin{aligned} \text{prox}_{sg^*}(y) &= y - s \text{prox}_{g/S}(y/s) \\ &= y - sb. \end{aligned}$$

This involves no linear systems!

Google, Nvidia, Gurobi and other big companies have implemented a solver known as PDLP based on this update (plus many other enhancements).

A primal-dual guarantee.

Our nonasymptotic convergence rate only ensured that

$$\min \|z_k - z_{k+1}\| = O\left(\frac{1}{\sqrt{k}}\right)$$

For primal-dual problems this might be a good progress metric.

Problems (P) - (D) can also be formulated as a minimax problem

$$\inf_x \sup_y f(x) + \langle Ax, y \rangle - h^*(y)$$

$\underbrace{\qquad\qquad\qquad}_{L(x, y)}$

Indeed a simple computation
yields primal-dual solution

$$L(x^*, y) \leq L(\overbrace{x^*, y^*}) \leq L(x, y^*)$$

Thus, a natural measure of
fitness is the primal-dual
gap

$$L(x, y^*) - L(x^*, y).$$